

Data Wrangling with dplyr


The R Bootcamp
Twitter: [@therbootcamp](https://twitter.com/therbootcamp)

September 2017

dplyr

dplyr is a package for managing dataframes.

Anytime you want to slice, dice, aggregate, or manipulate a dataframe, there is almost certainly a way to do it in dplyr.



Questions you can answer with dplyr

Can you calculate the mean survival times for each treatment separated by gender and time?

I need to know the mean birth rate only for countries in Africa from 1980 to 1980.

What percent of female patients had adverse events to drug X during weeks 5 through 10?

dplyr CheatSheet!

<https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

Data Wrangling with dplyr and tidy

Cheat Sheet

R Studio

Syntax - Helpful conventions for wrangling

dplyr::tbl_df(iris)
Converts data to tbl class. tbl's are easier to examine than data frames. R displays only the data that fits onscreen:

```
Source: local data frame [150 x 5]
  Sepal.Length Sepal.Width Petal.Length
1           5.1           3.5           1.4
2           4.9           3.0           1.4
3           4.7           3.2           1.3
4           4.6           3.1           1.5
5           5.0           3.6           1.4
..
Variables not shown: Petal.Width (dbl),
Species (fctr)
```

dplyr::glimpse(iris)
Information dense summary of tbl data.

utils::View(iris)
View data set in spreadsheet-like display (note capital V).

```
Source: local data frame [150 x 5]
  Sepal.Length Sepal.Width Petal.Length Species
1           5.1           3.5           1.4 setosa
2           4.9           3.0           1.4 setosa
3           4.7           3.2           1.3 setosa
4           4.6           3.1           1.5 setosa
5           5.0           3.6           1.4 setosa
6           5.4           3.7           1.4 setosa
7           4.6           3.4           1.4 setosa
8           5.0           3.4           1.5 setosa
```

dplyr::%>%
Passes object on left hand side as first argument (or argument) of function on righthand side.

$x \%>\% f(y)$ is the same as $f(x, y)$
 $y \%>\% f(x, \dots, z)$ is the same as $f(x, y, z)$

"Piping" with %>% makes code more readable, e.g.

```
iris %>%
  group_by(Species) %>%
  summarise(avg = mean(Sepal.Length)) %>%
  arrange(avg)
```

Tidy Data - A foundation for wrangling in R

In a tidy data set:

- Each variable is saved in its own column
- Each observation is saved in its own row

Tidy data complements R's **vectorized operations**. R will automatically preserve observations as you manipulate variables. No other format works as intuitively with R.

$M * A \rightarrow F$

Reshaping Data - Change the layout of a data set

tidyr::gather(cases, "year", "n", 2:4)
Gather columns into rows.

tidyr::spread(pollution, size, amount)
Spread rows into columns.

tidyr::separate(storms, date, c("y", "m", "d"))
Separate one column into several.

tidyr::unite(data, col, ..., sep)
Unite several columns into one.

dplyr::data_frame(a = 1:3, b = 4:6)
Combine vectors into data frame (optimized).

dplyr::arrange(mtcars, mpg)
Order rows by values of a column (low to high).

dplyr::arrange(mtcars, desc(mpg))
Order rows by values of a column (high to low).

dplyr::rename(tb, y = year)
Rename the columns of a data frame.

Subset Observations (Rows)

dplyr::filter(iris, Sepal.Length > 7)
Extract rows that meet logical criteria.

dplyr::distinct(iris)
Remove duplicate rows.

dplyr::sample_frac(iris, 0.5, replace = TRUE)
Randomly select fraction of rows.

dplyr::sample_n(iris, 10, replace = TRUE)
Randomly select n rows.

dplyr::slice(iris, 10:15)
Select rows by position.

dplyr::top_n(storms, 2, date)
Select and order top n entries (by group if grouped data).

Subset Variables (Columns)

dplyr::select(iris, Sepal.Width, Petal.Length, Species)
Select columns by name or helper function.

Helper functions for select - ?select

```
select(iris, contains("l"))
Select columns whose name contains a character string.
select(iris, ends_with("Length"))
Select columns whose name ends with a character string.
select(iris, everything())
Select every column.
select(iris, matches("l.*"))
Select columns whose name matches a regular expression.
select(iris, num_range("x", 1:5))
Select columns named x1, x2, x3, x4, x5.
select(iris, one_of(c("Species", "Genus")))
Select columns whose names are in a group of names.
select(iris, starts_with("Sepal"))
Select columns whose name starts with a character string.
select(iris, Sepal.Length:Petal.Width)
Select all columns between Sepal.Length and Petal.Width (inclusive).
select(iris, -Species)
Select all columns except Species.
```

Logic in R - ?Comparison, ?base::Logic			
<	Less than	!=	Not equal to
>	Greater than	%in%	Group membership
==	Equal to	is.na	Is NA
<=	Less than or equal to	!is.na	Is not NA
>=	Greater than or equal to	%, , !, xof, any, all	Boolean operators

devtools::install_github("rstudio/EDAWR") for data sets


Learn more with `browseVignettes(package = c("dplyr", "tidyr"))` • dplyr 0.4.0 • tidyr 0.2.0 • Updated: 1/15

dplyr

dplyr is a combination of 3 things:

1. **objects** like dataframes
2. **verbs** that **do** things to objects.
3. **pipes** `%>%` that string together objects and verbs

Sequential



dplyr is meant to be sequential and work like language

Take data X, then do Y, then do Z...

Here's the basic structure of dplyr in action

```
data %>%           # Start with data, and THEN  
  VERB1 %>%        # Do VERB1, (and THEN)  
  VERB2 %>% ...    # Do VERB2, (and THEN)
```

dplyr

Question:

From the ChickWeight dataframe, calculate the mean weight and time for each diet

Answer:

```
library(dplyr)

x <- ChickWeight %>% # Start with ChickWeight
  group_by(Diet) %>% # Group by Diet
  summarise(         # Get ready to summarise....
    weight.mean = mean(weight), # Mean weight
    time.mean = mean(Time),     # Mean time
    N = n()                    # Number of cases
  )
```

x

```
## # A tibble: 4 x 4
##   Diet weight.mean time.mean     N
##   <fctr>      <dbl>    <dbl> <int>
## 1     1      102.6     10.48   220
## 2     2      122.6     10.92   120
## 3     3      142.9     10.92   120
## 4     4      135.3     10.75   118
```

Common dplyr verbs

verb	action	example
<code>filter()</code>	Select rows based on some criteria	<code>filter(age > 40 & sex == "m")</code>
<code>arrange()</code>	Sort rows	<code>arrange(date, group)</code>
<code>select()</code>	Select columns (and ignore all others)	<code>select(age, sex)</code>
<code>rename()</code>	Rename columns	<code>rename(DATE_MONTHS_X24, date)</code>
<code>mutate()</code>	Add new columns	<code>mutate(height.m = height.cm / 100)</code>
<code>case_when()</code>	Recode values of a column	<code>sex.n = case_when(sex == 0 ~ "m", sex == 1 ~ "f")</code>
<code>group_by()</code> , <code>summarise()</code>	Group data and then calculate summary statistics	<code>group_by(treatment) %>% summarise(...)</code>

Example 1

Add a column called `weight_d_time` that is weight divided by time

```
library(dplyr)

x <- ChickWeight %>%           # Start with the ChickWeight data
  mutate(                       # Create new columns...
    weight_d_time = weight / Time
  )

head(x)    # Print the result
```

##	weight	Time	Chick	Diet	weight_d_time
## 1	42	0	1	1	Inf
## 2	51	2	1	1	25.50
## 3	59	4	1	1	14.75
## 4	64	6	1	1	10.67
## 5	76	8	1	1	9.50
## 6	93	10	1	1	9.30

Example 2

Add a column called *weight_d_time* that is weight divided by time AND *time_d* that is time in days

```
x <- ChickWeight %>%           # Start with the ChickWeight data
  mutate(                      # Create new columns...
    weight_d_time = weight / Time, # weight_d_time is weight divided by Time
    time_d = Time * 7             # time_d is Time times 7
  )

head(x) # Print the result
```

##	weight	Time	Chick	Diet	weight_d_time	time_d
## 1	42	0	1	1	Inf	0
## 2	51	2	1	1	25.50	14
## 3	59	4	1	1	14.75	28
## 4	64	6	1	1	10.67	42
## 5	76	8	1	1	9.50	56
## 6	93	10	1	1	9.30	70

Recoding values with case_when()

Recoding values is a common data wrangling task. You can easily do this with `case_when()`:

```
data %>%  
  mutate(  
    var_new = case_when(  
      var_old == OLD_A ~ NEW_A,  
      var_old == OLD_B ~ NEW_B  
    )  
  )
```

For example, in a dataset, the column `sex` might be coded with 1s and 0s.

You might want to create a new column `sex_new` where 1 = "female" and 0 = "male":

sex	sex_new
1	"female"
0	"male"

To change the value of 1 to "female", and 0 to "male", you can use `case_when()`:

```
# Add a column sex_new to data  
  
data <- data %>%  
  mutate(  
    sex_new = case_when(  
      sex == 1 ~ "female",  
      sex == 0 ~ "male"  
    )  
  )
```

You can think about the code above as follows:

- Create a new column `sex_new` where
 - If `sex == 1`, then set the value to "female"
 - If `sex == 0`, then set the value to "male"

Example 3

Create a new variable `Diet_name` which shows `Diet` in text format. Here is a table of the values

Diet	Diet_name
1	"fruit"
2	"vegetables"
3	"meat"
4	"grains"

```
ChickWeight <- ChickWeight %>% # Start
  mutate(
    Diet_name = case_when(
      Diet == 1 ~ "fruit",
      Diet == 2 ~ "vegetables",
      Diet == 3 ~ "meat",
      Diet == 4 ~ "grains"
    )
  )
head(ChickWeight)
```

```
##   weight Time Chick Diet Diet_name
## 1     42    0     1   1    fruit
## 2     51    2     1   1    fruit
## 3     59    4     1   1    fruit
## 4     64    6     1   1    fruit
## 5     76    8     1   1    fruit
## 6     93   10     1   1    fruit
```

Example 4

For each Diet, calculate the mean weight

```
ChickWeight %>%           # Start with the ChickWeight data
  group_by(Diet) %>%       # Group the data by Diet
  summarise(               # Now summarise....
    weight.mean = mean(weight) # Mean weight
  )
```

```
## # A tibble: 4 x 2
##   Diet weight.mean
##   <fctr>      <dbl>
## 1     1      102.6
## 2     2      122.6
## 3     3      142.9
## 4     4      135.3
```

Example 5

For each time period less than 10, calculate the mean weight

```
ChickWeight %>%           # Start with the ChickWeight data
  filter(Time < 10) %>%    # Only Time periods less than 10
  group_by(Time) %>%      # Group the data by Diet
  summarise(              # Now summarise....
    weight.mean = mean(weight) # Mean weight
  )
```

```
## # A tibble: 5 x 2
##   Time weight.mean
##   <dbl>      <dbl>
## 1     0      41.06
## 2     2      49.22
## 3     4      59.96
## 4     6      74.31
## 5     8      91.24
```

Example 6

For each Diet, calculate the mean weight, maximum time, and the number of chicks on each diet:

```
ChickWeight %>%           # Start with the ChickWeight data
  group_by(Diet) %>%       # Group the data by Diet
  summarise(               # Now summarise....
    weight.mean = mean(weight), # Mean weight
    time.max = max(Time),      # Max time
    N = n()                   # Number of observations
  )
```

```
## # A tibble: 4 x 4
##   Diet weight.mean time.max    N
##   <fctr>      <dbl>    <dbl> <int>
## 1     1      102.6      21    220
## 2     2      122.6      21    120
## 3     3      142.9      21    120
## 4     4      135.3      21    118
```

Other dplyr verbs

verb	action	example
<code>sample_n()</code>	Select a random sample of n rows	<code>sample_n(10)</code>
<code>sample_frac()</code>	Select a random fraction of rows	<code>sample_frac(.20)</code>
<code>first()</code> , <code>last()</code>	Give the first (or last) observation	<code>first()</code> , <code>last()</code>

Example 7

Give me a random sample of 10 rows from the ChickWeight dataframe, but only show me the values for Chick and weight

```
# Give me a random sample of 10 rows, but only show me columns Chick and weight
```

```
ChickWeight %>%  
  select(Chick, weight) %>%  
  sample_n(10)
```

```
##      Chick weight  
## 46      4     154  
## 343     31      62  
## 61      6      41  
## 429     38      98  
## 113     10      81  
## 466     41     124  
## 547     48     104  
## 482     42     234  
## 335     30     115  
## 33      3     163
```

dplyr

dplyr operations (almost) always return a dataframe which you can assign to a new object:

Create a dataframe with the average weight for each time period and nothing else!!

```
# Create a new object called time_agg
```


```
time_agg <- ChickWeight %>%  
  group_by(Time) %>%  
  summarise(  
    weight.mean = mean(weight)  
  )  
  
head(time_agg)
```

```
## # A tibble: 6 x 2  
##   Time weight.mean  
##   <dbl>         <dbl>  
## 1     0         41.06  
## 2     2         49.22  
## 3     4         59.96  
## 4     6         74.31  
## 5     8         91.24  
## 6    10        107.84
```


dplyr summary

dplyr is great for elegantly performing sequential operations on data.

The 'pipe' operator `%>%` helps you string multiple *objects* (like dataframes) and *verbs* (summarise, order, aggregate...) together.



Basic structure of dplyr commands:

```
data %>%      # Start with data, AND THEN...
  VERB1 %>%   # Do VERB1, AND THEN...
  VERB2 %>%   # Do VERB2, AND THEN...
  VERB3 %>%   # Do VERB3, AND THEN...
  group_by(x, y) %>% # Group by variables x, y
    summarise(
      VAR_A_New = fun(X),
      VAR_B_New = fun(Y)
    )
)
```

Questions?

Wrangling Pratical

[Link to Wrangling practical](#)

<!--

-->

<!--

-->

<!--

-->

<!--

-->